



Deep learning: an overview of datasets and algorithms

DL taskforce Caterpillar

Rob Romijnders

Eindhoven, University of Technology

RomijndersRob@gmail.com

May 25, 2017

Overview

1 Why

2 How

3 Datasets

4 What

5 Wrap up

6 Questions

Object detection

Faster R-CNN

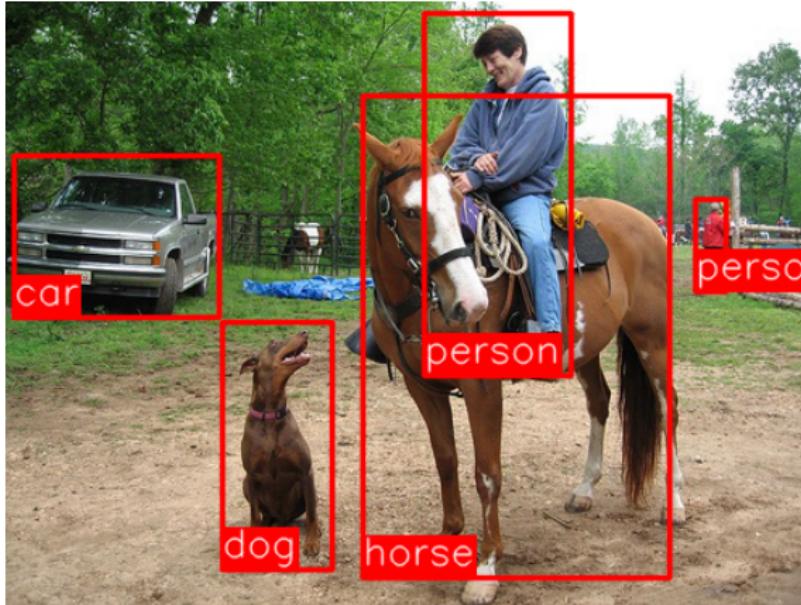


Figure: faster rcnn (image from Github *mitmul*)

Apple siri



Figure: Photo: cultofmac.com

GMail reply

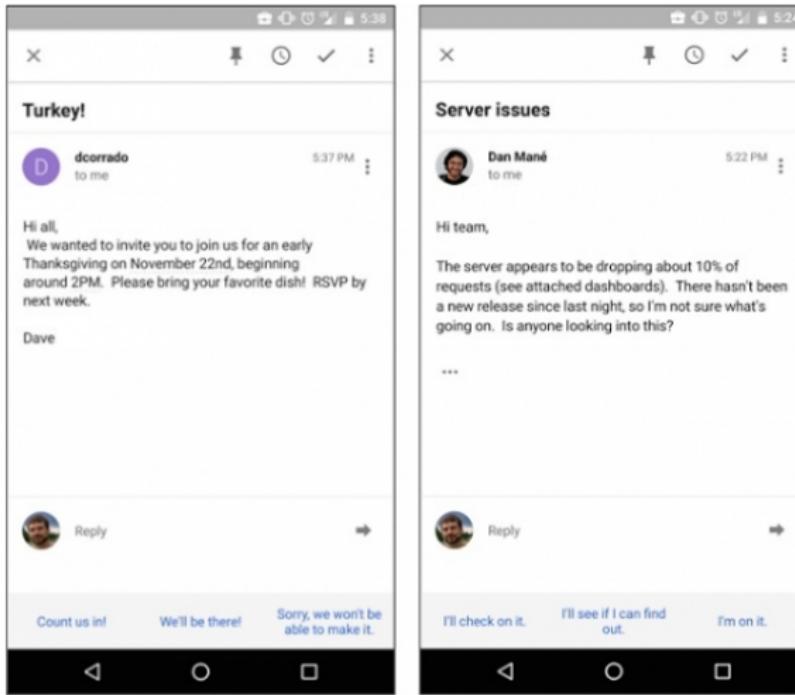
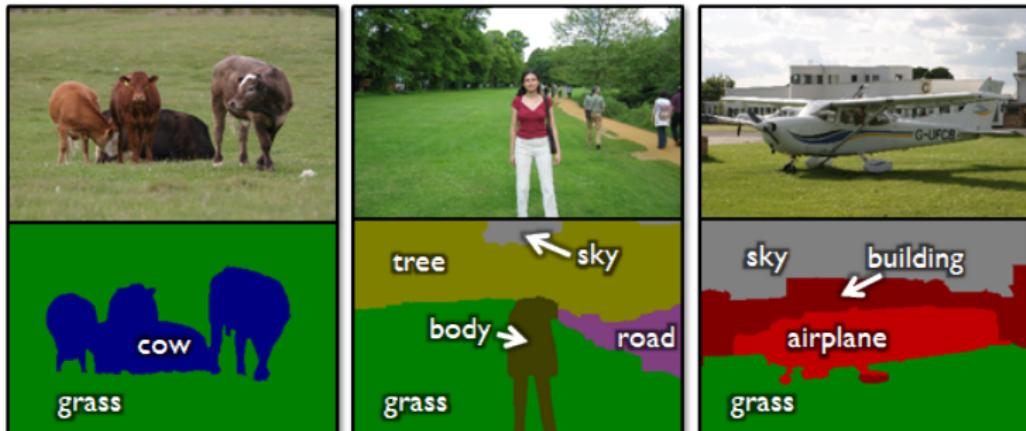


Figure: Photo: Greg Corrado, Google Research Blog

Segmentation



object classes	building	grass	tree	cow	sheep	sky	airplane	water	face	car
bicycle	flower	sign	bird	book	chair	road	cat	dog	body	boat

Figure: Semantic Segmentation with CNN (image from jamie.shotton.org)

Generation

THE MULTIVERSE —

Movie written by algorithm turns out to be hilarious and intense

For *Sunspring*'s exclusive debut on Ars, we talked to the filmmakers about collaborating with an AI.

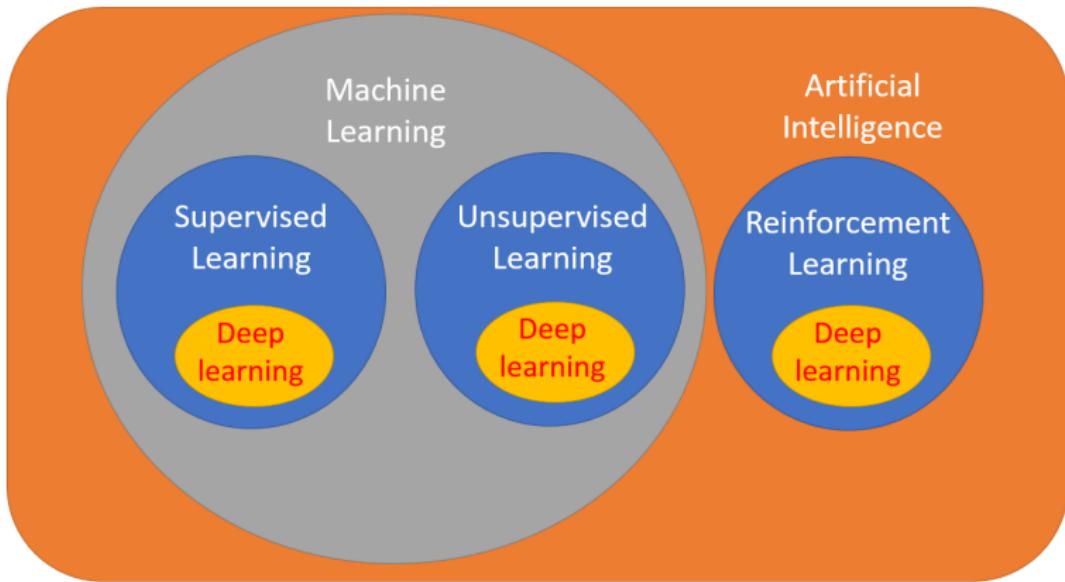
ANNALEE NEWITZ - 6/9/2016, 12:30 PM

Sunspring, a short science fiction movie written entirely by AI, debuts exclusively on Ars today.

Figure: <http://arstechnica.com/the-multiverse/2016/06/an-ai-wrote-this-movie-and-its-strangely-moving/>

Why not

Overview of the field



Which data not

- ① YES: media type data
 - ① Text, language, speech
 - ② Images, video, maps
 - ③ time-series, stocks, valuta
- ② NO: categorical data
 - ① Properties of instances
 - ② Features of instances
 - ③ Categories of products

Neural Networks

Neural nets

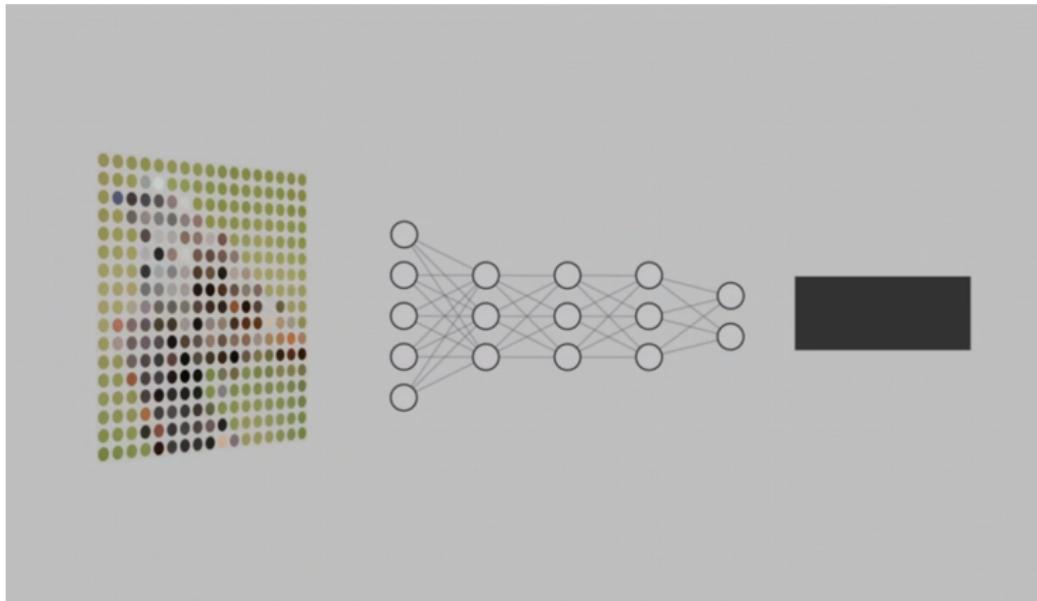


Figure: Neural network explained (credits: Blaise Aguera y Arcas)

Basic equation

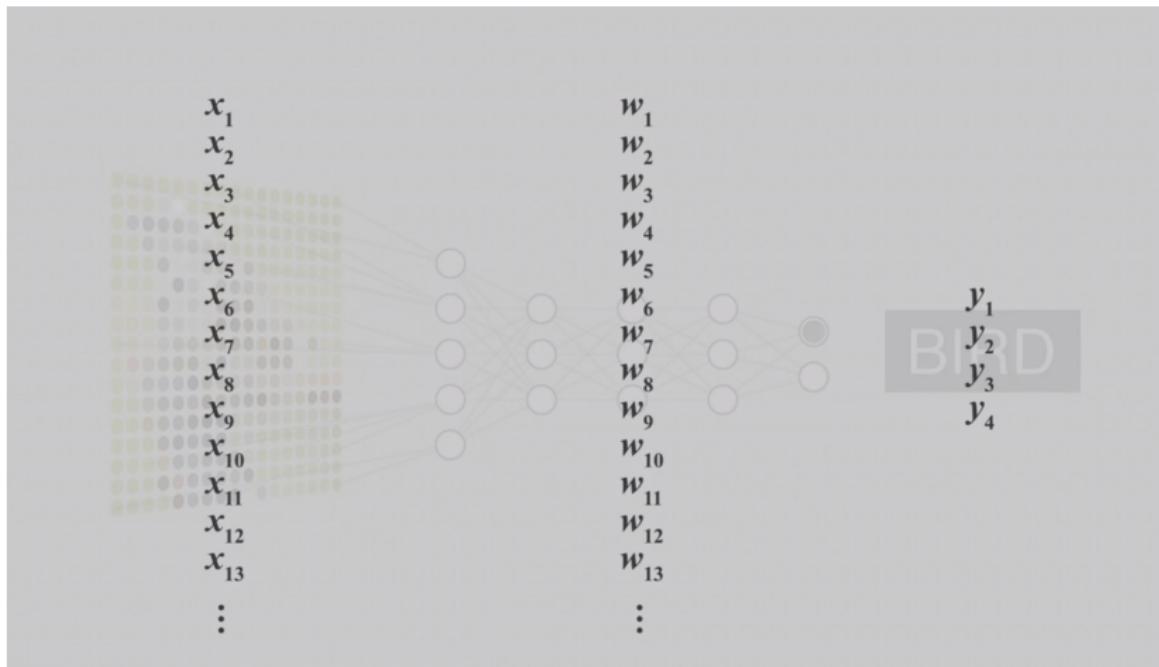


Figure: Neural net in algebraic form (credits: Blaise Aguera y Arcas)

Template equation



Figure: Template equation neural net (credits: Blaise Aguera y Arcas)



How to use

How to use them

$$w \cdot x = y$$

$$2 \cdot 3 = y$$

Forward inference

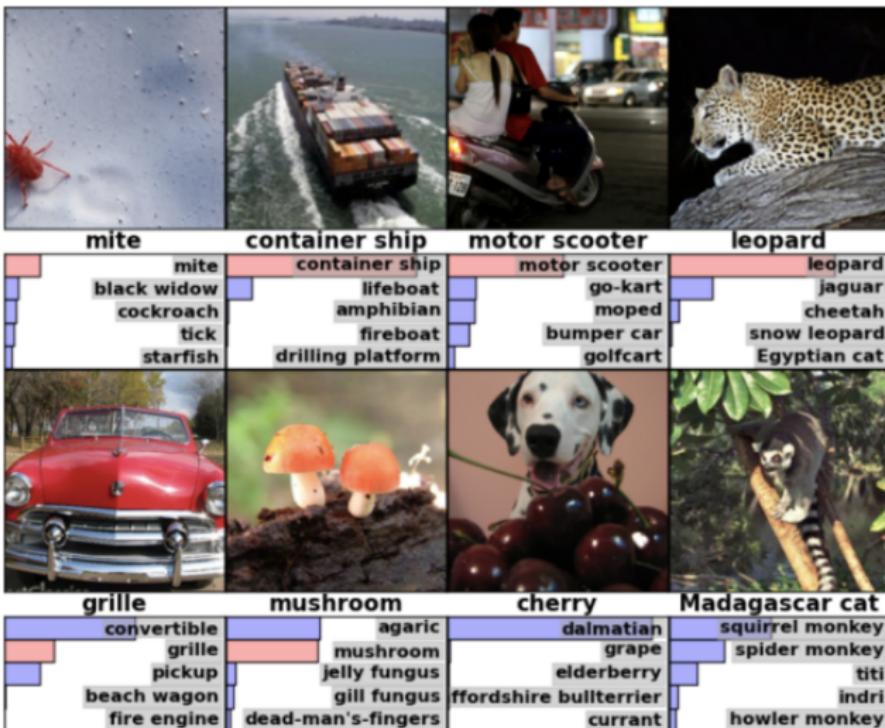


Figure: Forward inference CNN (Krizhevsky et al. 2012)

How to train

How to train them

$$w \cdot x = y$$

$$w \cdot 3 = 6$$

$$y \div x = w$$



$$\begin{aligned}0 &= w \cdot x - y \\0 &= w \cdot 3 - 6\end{aligned}$$

$$\text{error} = |w \cdot x - y| \rightarrow 0$$

Loss function

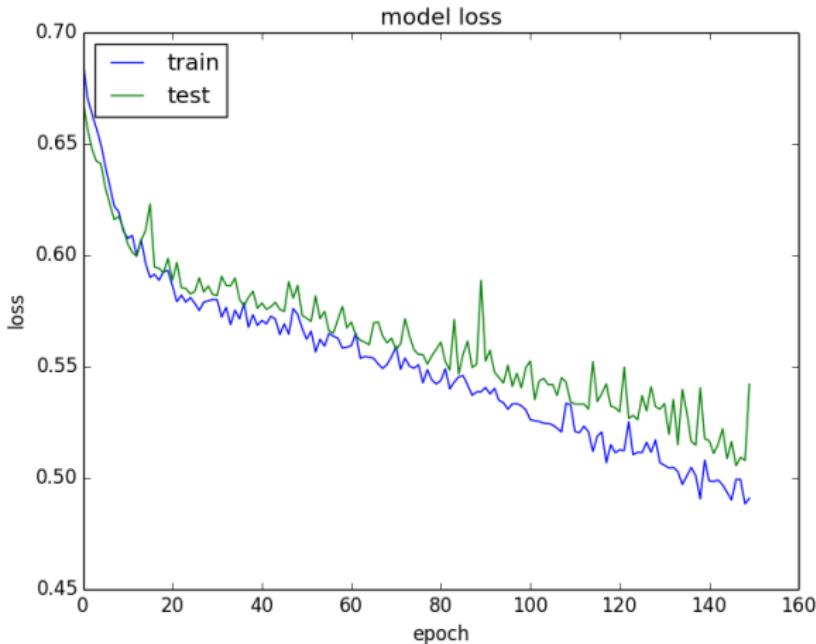


Figure: (<http://machinelearningmastery.com/>)

Architecture

Architecture

Architecture

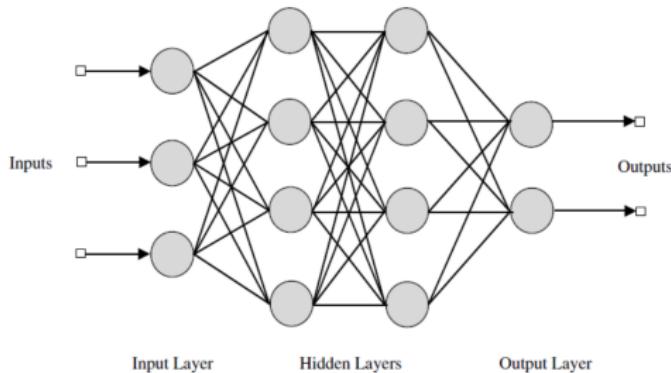


Figure: Feedforward neural network

$$y = w \cdot x$$

Two architectures

① Convolution

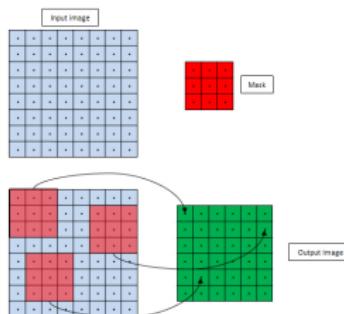


Figure: developer.amd.com

② Recursion

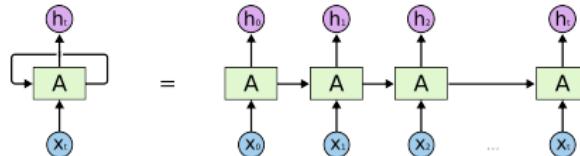


Figure: credits: colah.github.io

Convolution

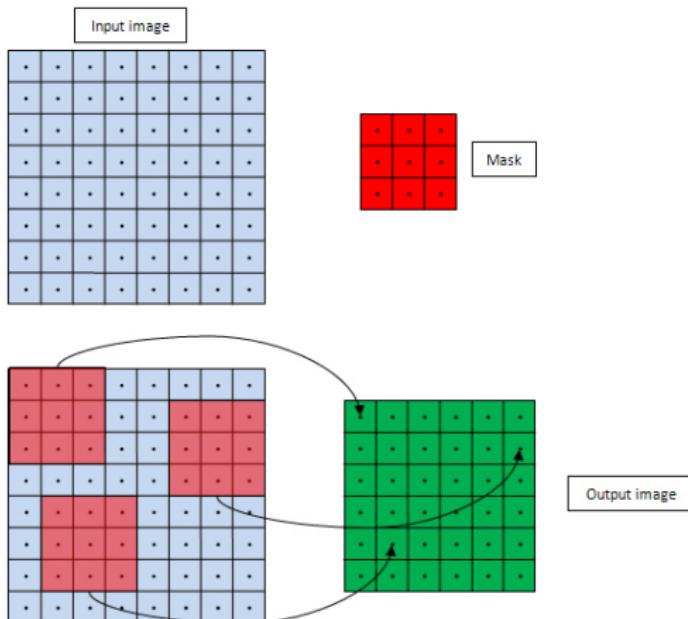


Figure: developer.amd.com

CNN

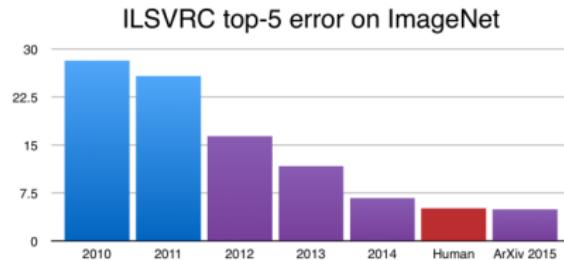


Figure: ImageNet results (credits jackkelly.github.io/)



Figure: Forward inference CNN (Krizhevsky et al. 2012)

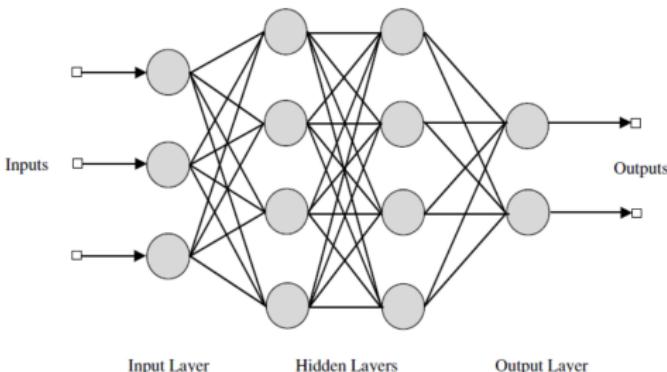
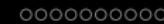
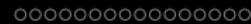


Figure: Feedforward neural network

$$32 \cdot 32 \cdot 100 = 102400$$

$$5 \cdot 5 = 25$$



Recursion

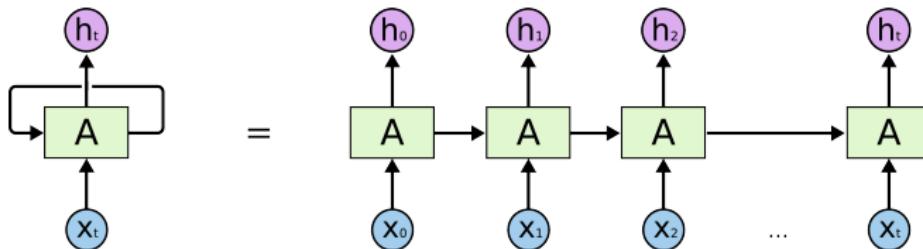


Figure: credits: colah.github.io



Template formula

$$y = w \cdot x$$

$$y_t = w \cdot (x_t, h_{t-1})$$



Formula

$$h = \phi(W_{xh}x)$$

$$y = \phi(W_{hy}h)$$

$$h_t = \phi(W_{xh}x_t + W_{hh}h_{t-1}) \quad y_t = \phi(W_{hy}h_t)$$



Python code

```
rnn = RNN()  
y = rnn.step(x)
```

```
class RNN:  
    def step(self, x):  
        self.h = np.tanh(np.dot(self.Whh, self.h)  
                        + np.dot(self.Wxh, x))  
        y = np.dot(self.Why, self.h)  
    return y
```

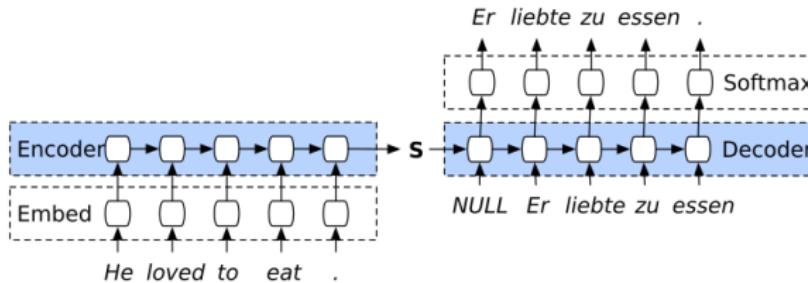


Figure: RNN for Machine Translation (credits: smerity.com)



Example

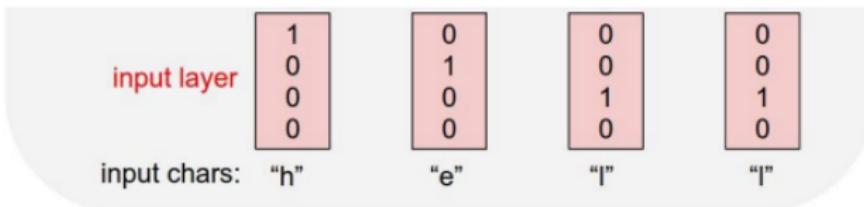


Figure: cs231n, Andrej Karpathy

Example

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

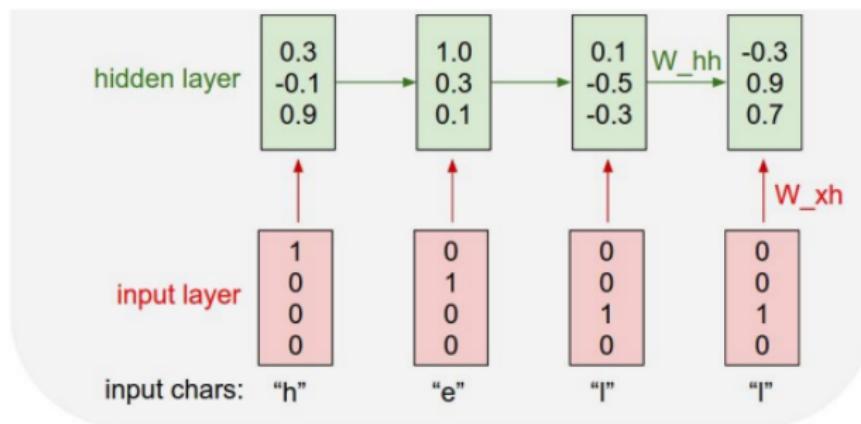


Figure: cs231n, Andrej Karpathy

Example

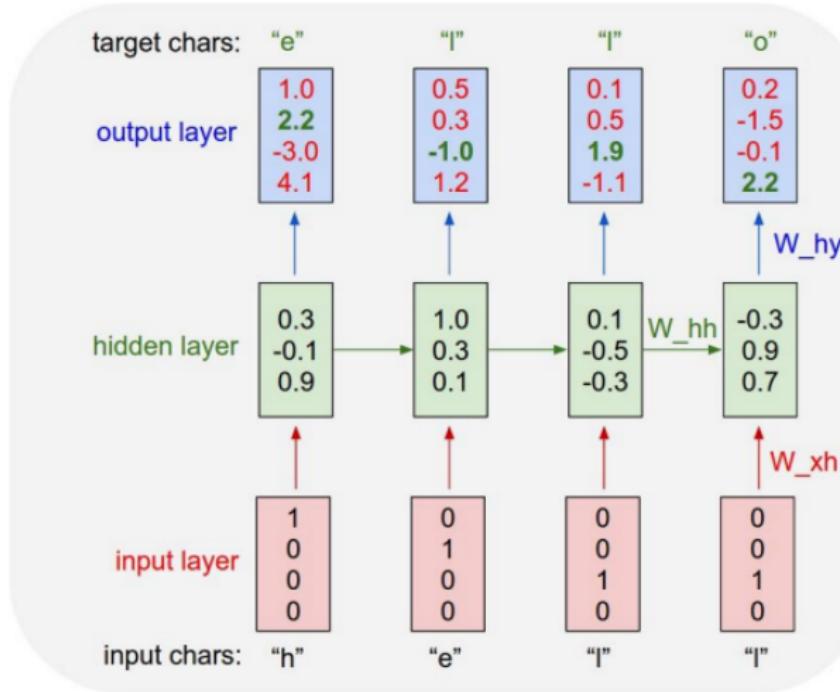


Figure: cs231n, Andrej Karpathy



Long Short-term memory

$$i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1} + b_i)$$

$$f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1} + b_f)$$

$$o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1} + b_o)$$

$$c_t = f_t c_{t-1} + i_t \tanh(W_{xc}x_t + W_{hc}h_{t-1} + b_c)$$

$$h_t = o_t \tanh(c_t)$$

LSTM

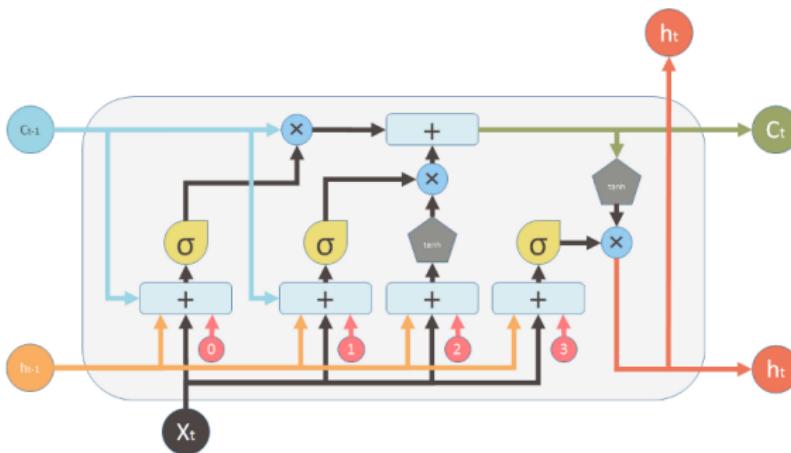


Figure: Diagram depicting LSTM block

Author: Shi Yan. Source: medium.com/@oshiyan/

Overview			
	FF	RNN	CNN
- Basic computation	Matrix multiplies	Recurrent multiplications	Convolutions
- Typical data	Embedding vectors	Sequences	Images
- Weights ties	No	Sequential	Spatial
- Effective span	Across input	Long term	Local receptive field
- Stationary		Seldom events	Frequent patterns
- How to train	Backprop	Backprop	Backprop
- Performance	No guarantee	No guarantee	No guarantee
- Typical examples	None	Machine translation, speech tagging, text generation	segmentation, image processing, object detection

Span

① CNN

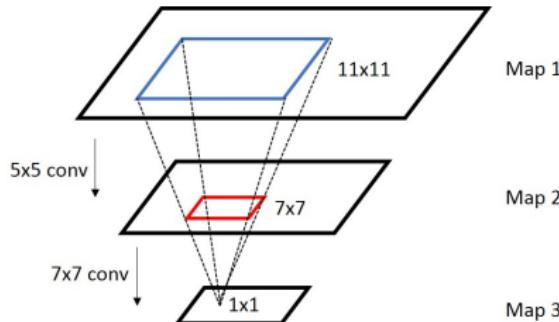


Figure: credits: cvmarcher.com/

② RNN

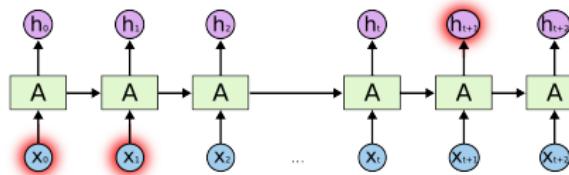


Figure: credits colah.github.io

Datasets and algorithms

Tagging of clinical events

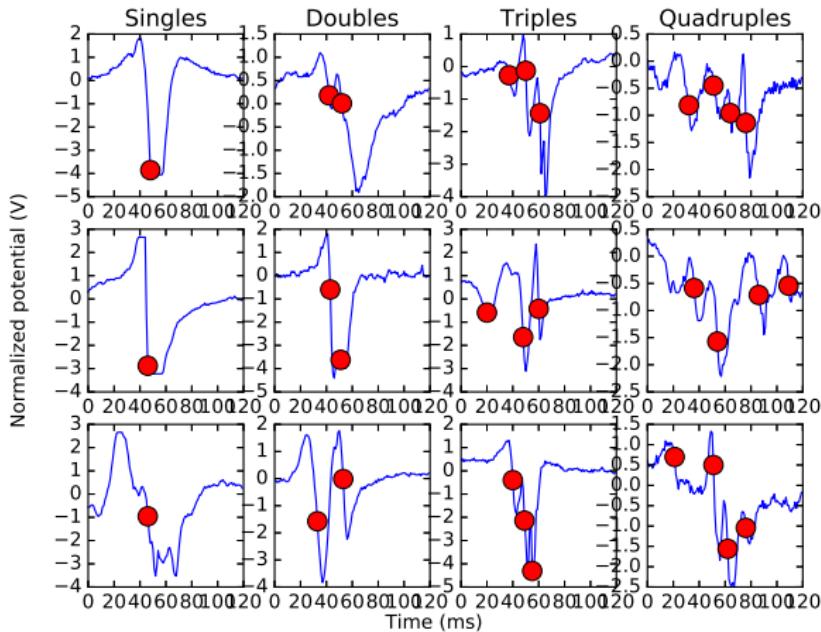


Figure: Classification of fractionated electrograms in epicardial mappings using a recurrent neural network, R Romijnders et al.

Use LSTM's, because the deflections occur as events

Classifying music

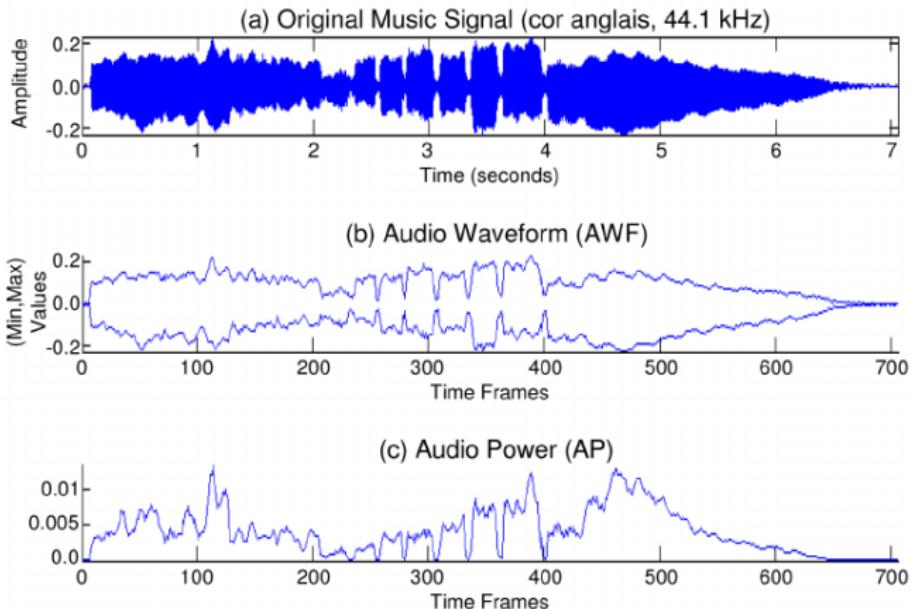


Figure: ML for music



Use CNN's, because music genre follows a certain pattern



- ① *The film did n't move me one way or the other , but it was an honest effort and if you want to see a flick about telemarketers this one will due .*
- ② *For those of us who respond more strongly to storytelling than computer-generated effects , the new Star Wars installment has n't escaped the rut dug by the last one .*

IMDB sentiment challenge, Kaggle



Use LSTM, because a single word could change the sentiment. Compare *I do like this* and *I do not like this*

Architectures for RNN

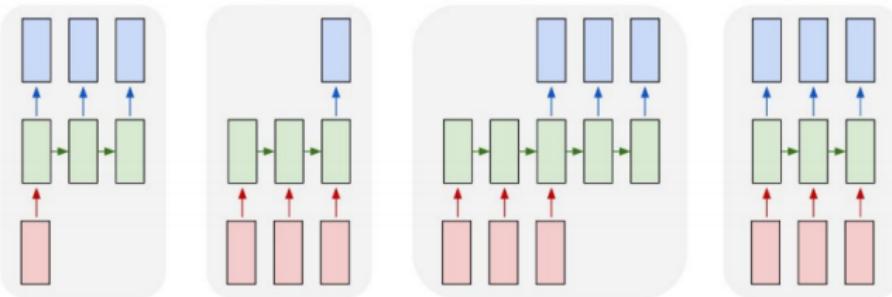


Figure: credits: [Andrej Karpathy](#)

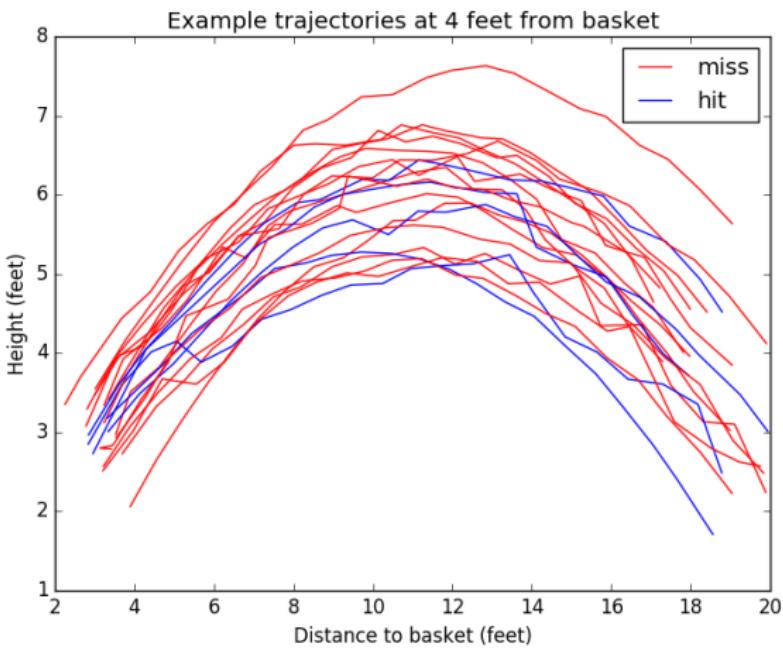
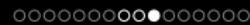


Figure: DL for basketball, Rajiv and Rob



Works with both CNN and LSTM



PANDARUS:

Alas, I think he shall be come approached and the day
When little strain would be attain'd into being never fed,
And who is but a chain and subjects of his death,
I should not sleep.

Second Senator:

They are away this miseries, produced upon my soul,
Breaking and strongly should be buried, when I perish
The earth and thoughts of many states.

DUKE VINCENTIO:

Well, your wit is in the care of side and that.

Second Lord:

They would be ruled after this chamber, and
my fair nues begun out of the fact, to be conveyed,
Whose noble souls I'll have the heart of the wars.

Clown:

Come, sir, I will make did behold your worship.

VIOLA:

I'll drink it.

Figure: credits: [Andrej Karpathy](#)



Use LSTM's, because it
predicts a vector (word) for
each time step



Tabular

Data Dictionary

Variable	Definition	Key
survival	Survival	0 = No, 1 = Yes
pclass	Ticket class	1 = 1st, 2 = 2nd, 3 = 3rd
sex	Sex	
Age	Age in years	
sibsp	# of siblings / spouses aboard the Titanic	
parch	# of parents / children aboard the Titanic	
ticket	Ticket number	
fare	Passenger fare	
cabin	Cabin number	
embarked	Port of Embarkation	C = Cherbourg, Q = Queenstown, S = Southampton

Figure: [Titanic KAGGLE](#)



Do NOT use deep learning. Go
for Random Forests or SVM



Captioning



Figure: [Link to page](#)



Use CNN for pattern extraction
from image. Use RNN to
generate caption.

DIY

- ➊ Gather data
- ➋ Choose architecture
- ➌ Code up neural net
- ➍ Train
- ➎ Add bells and whistles
- ➏ Post processing

Gather data

① SMALL:

- ① MNIST: 60.000 samples
- ② TIMIT: 630 speakers, 10 sentences

② NORMAL:

- ① Imagenet: 3.2 million images
- ② FLICKR: 1 million images
- ③ Word vectors on all of Wikipedia



Amazon Mechanical Turk



amazon mechanical turk

beta.

Choose architecture

① Convolution

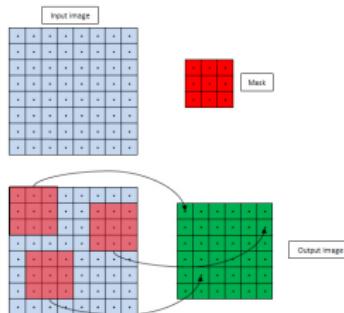


Figure: developer.amd.com

② Recursion

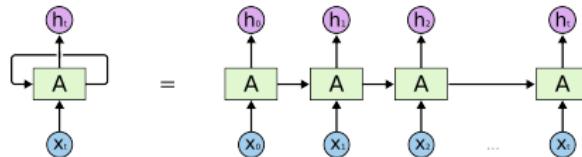


Figure: credits: colah.github.io

Code

① Low level

- ① Tensorflow
- ② Torch, pyTorch

② High level

- ① Keras
- ② Scikit flow, pretty tensor, tfslim, ...



Train

- ① No bells and whistles
- ② Small sub dataset
- ③ See if you can overfit

Bells and whistles

- ➊ Batch normalization, layer normalization
- ➋ Dropout
- ➌ Add more data: semi supervised learning
- ➍ Residual connections
- ➎ etcetera, etcetera



Activation functions

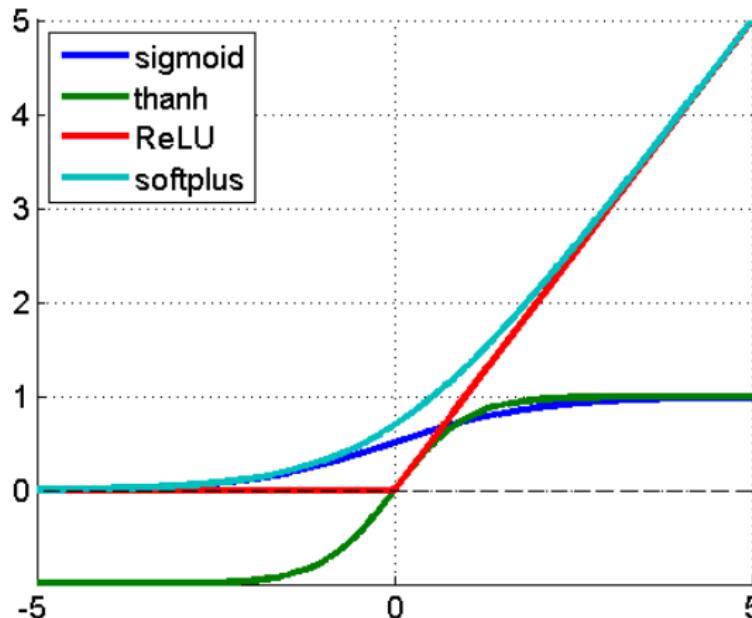


Figure: Common activation functions (credits Vanessa Imiloa)



Regularization



Figure: Example for regularization



MNIST

Random Sampling of MNIST

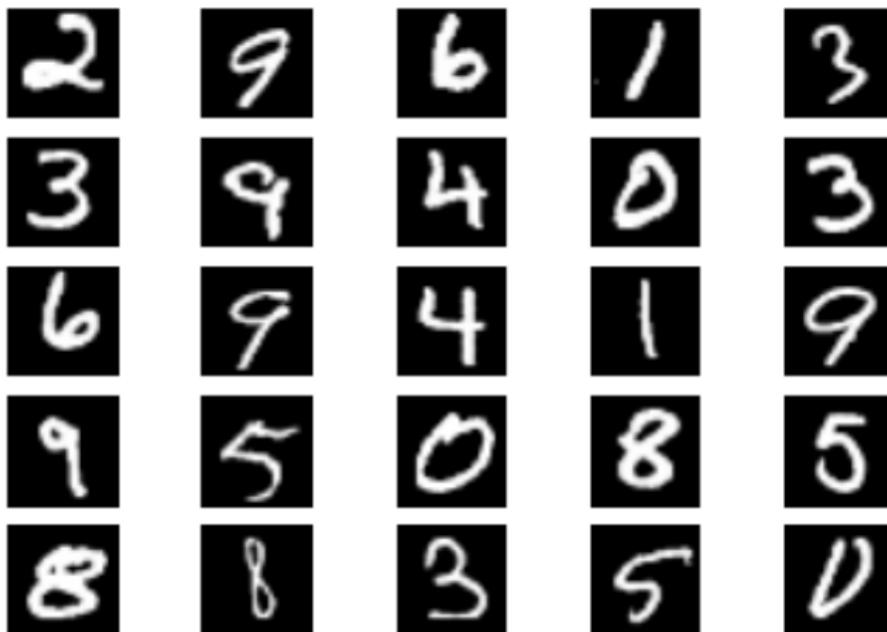


Figure: Caption



Post processing

- ① Plot low confidence / high uncertainty
- ② Plot confusion matrix
- ③ Inspect neurons

Two architectures

① Convolution

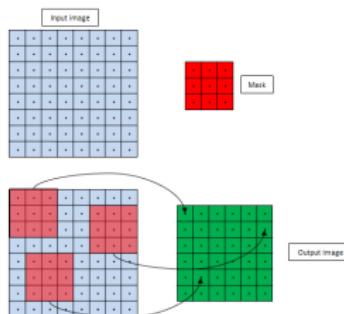


Figure: developer.amd.com

② Recursion

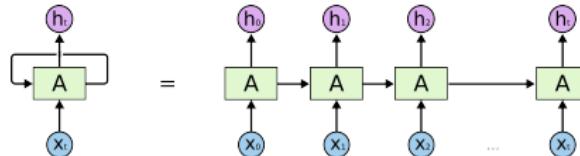


Figure: credits: colah.github.io

Questions



Further reading

① RNN

- ① The Unreasonable Effectiveness of Recurrent Neural Networks, Andrej Karpathy
- ② Supervised sequence labelling with RNN, Alex Graves
- ③ Colah's blog (colah.github.io)

② CNN

- ① Course on CNN: cs231n (cs231n.stanford.edu/syllabus.html)
- ② Udacity on Deep learning (udacity.com/course/deep-learning-ud730)
- ③ My page (robromijnders.github.io)